

Getting (your) scientific software installed using

Kenneth Hoste (HPC-UGent)

kenneth.hoste@ugent.be

20190510 - CSCS Software Management Course

https://users.ugent.be/~kehoste/EasyBuild_20190510_CSCS_software_management.pdf

<https://easybuilders.github.io/easybuild>

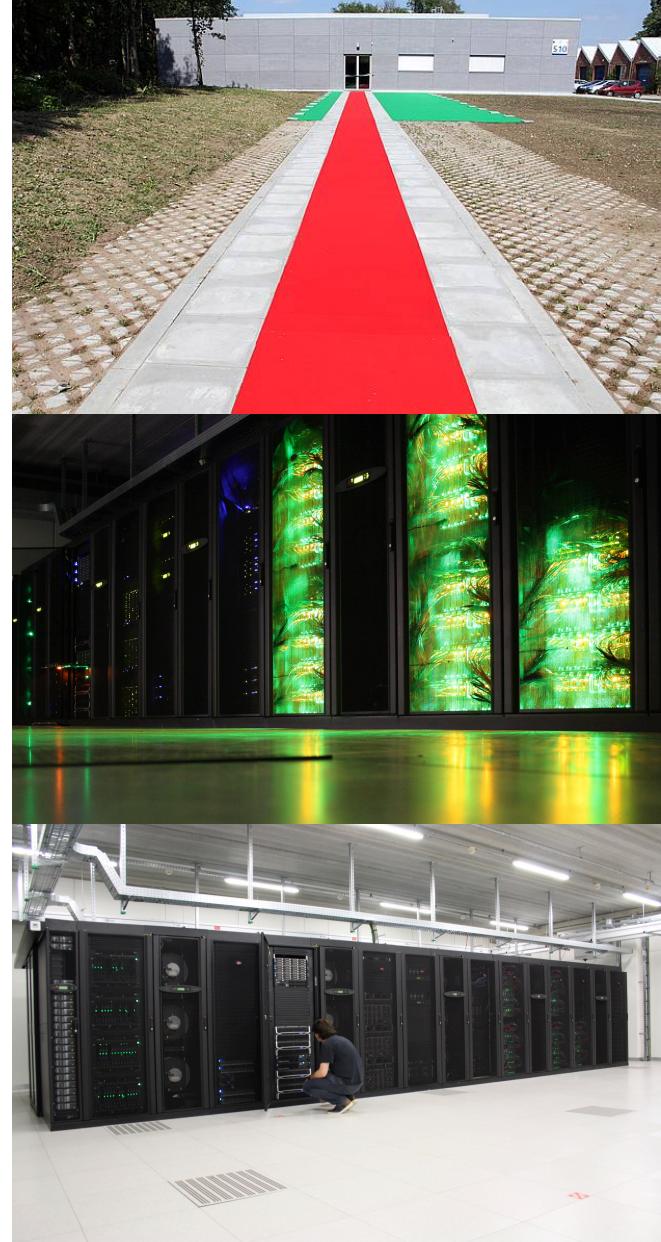
<https://easybuild.readthedocs.io>

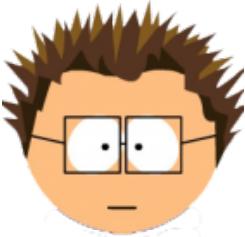
HPC-UGent



- part of central IT department of Ghent University (Belgium)
- centralised scientific computing services, training & support
- for researchers of UGent, industry & knowledge institutes
- 6 Tier-2 clusters (> 15k cores in total), ~2 PB shared storage
- 7+1 team members, > 3000 user accounts
- member of Flemish Supercomputer Centre (VSC)

<https://www.vscentrum.be>





whoami

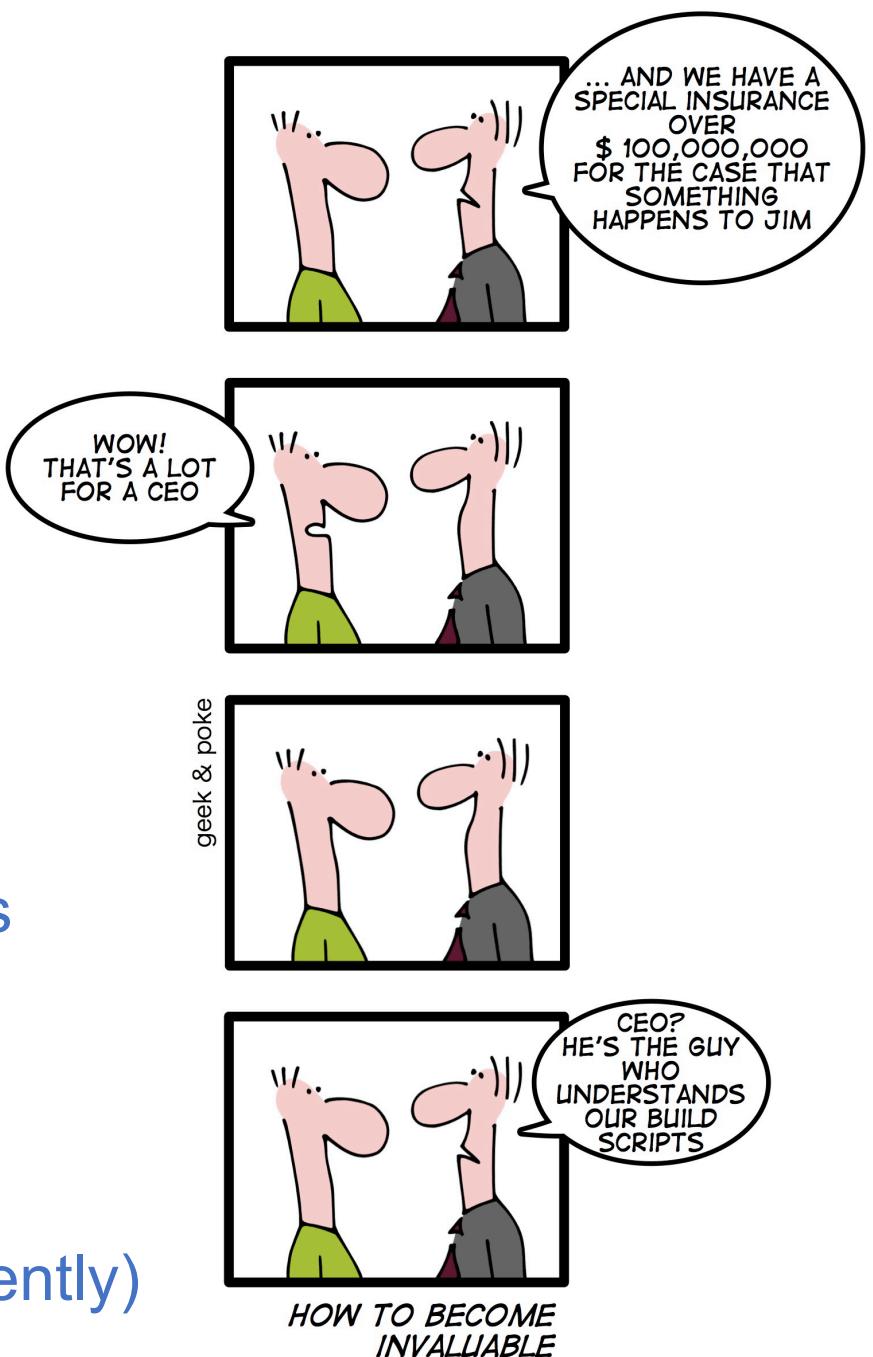
kenneth.hoste@ugent.be
[@boegel \(GitHub, IRC, Slack\)](https://github.com/boegel)
[@kehoste \(Twitter\)](https://twitter.com/kehoste)

- Masters & PhD in Computer Science from Ghent University
- PhD topic: machine learning applied to software performance, compilers, ...
- joined HPC-UGent team in October 2010
- main tasks: user support & training, *software installations*
- slowly also became  **EasyBuild lead developer & release manager**
- likes family, beer, loud music, FOSS, helping people, dad jokes, stickers, ...
- doesn't like CMake, SCons, Bazel, TensorFlow, OpenFOAM, ...

Getting scientific software installed

Installation of scientific software is a tremendous problem for HPC sites all around the world.

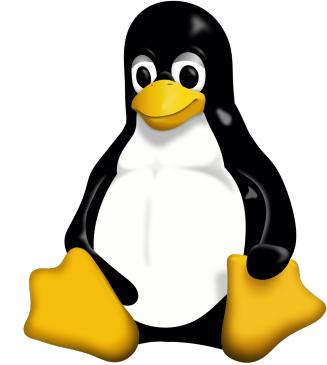
- ideally built from source (performance is key!)
- tedious, time-consuming, frustrating,
sometimes simply not worth the (manual) effort...
- huge burden on researchers & HPC support teams
 - over 25% of support tickets at HPC-UGent,
but consumes way more than 1/4th of support time...
- very little collaboration among HPC sites (until recently)



What about existing software installation tools? (1/2)

- **traditional package managers in Linux(-like) operating systems**

- *yum* to install RPMs (Red Hat derivatives)
- *apt* to install .deb files (Debian & derivatives)
- *Homebrew* (macOS/Linux)
- *Portage* (Linux), *pkgsrc* (*nix), ...



- **problems in context of scientific software & HPC:**

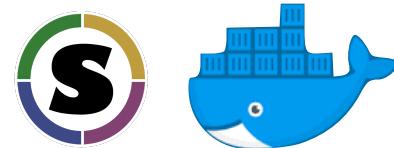
- not well suited to idiosyncrasies of scientific software
- not aware (enough) of MPI/BLAS/LAPACK/GPUs, other compilers (Intel, PGI, ...)
- bad fit for multi-user HPC systems in general
- little support for having multiple versions installed side-by-side
- strong focus on generically optimised binaries (portability is important to limit effort)

What about existing software installation tools? (2/2)

CONDA (installation tool for Anaconda distribution)

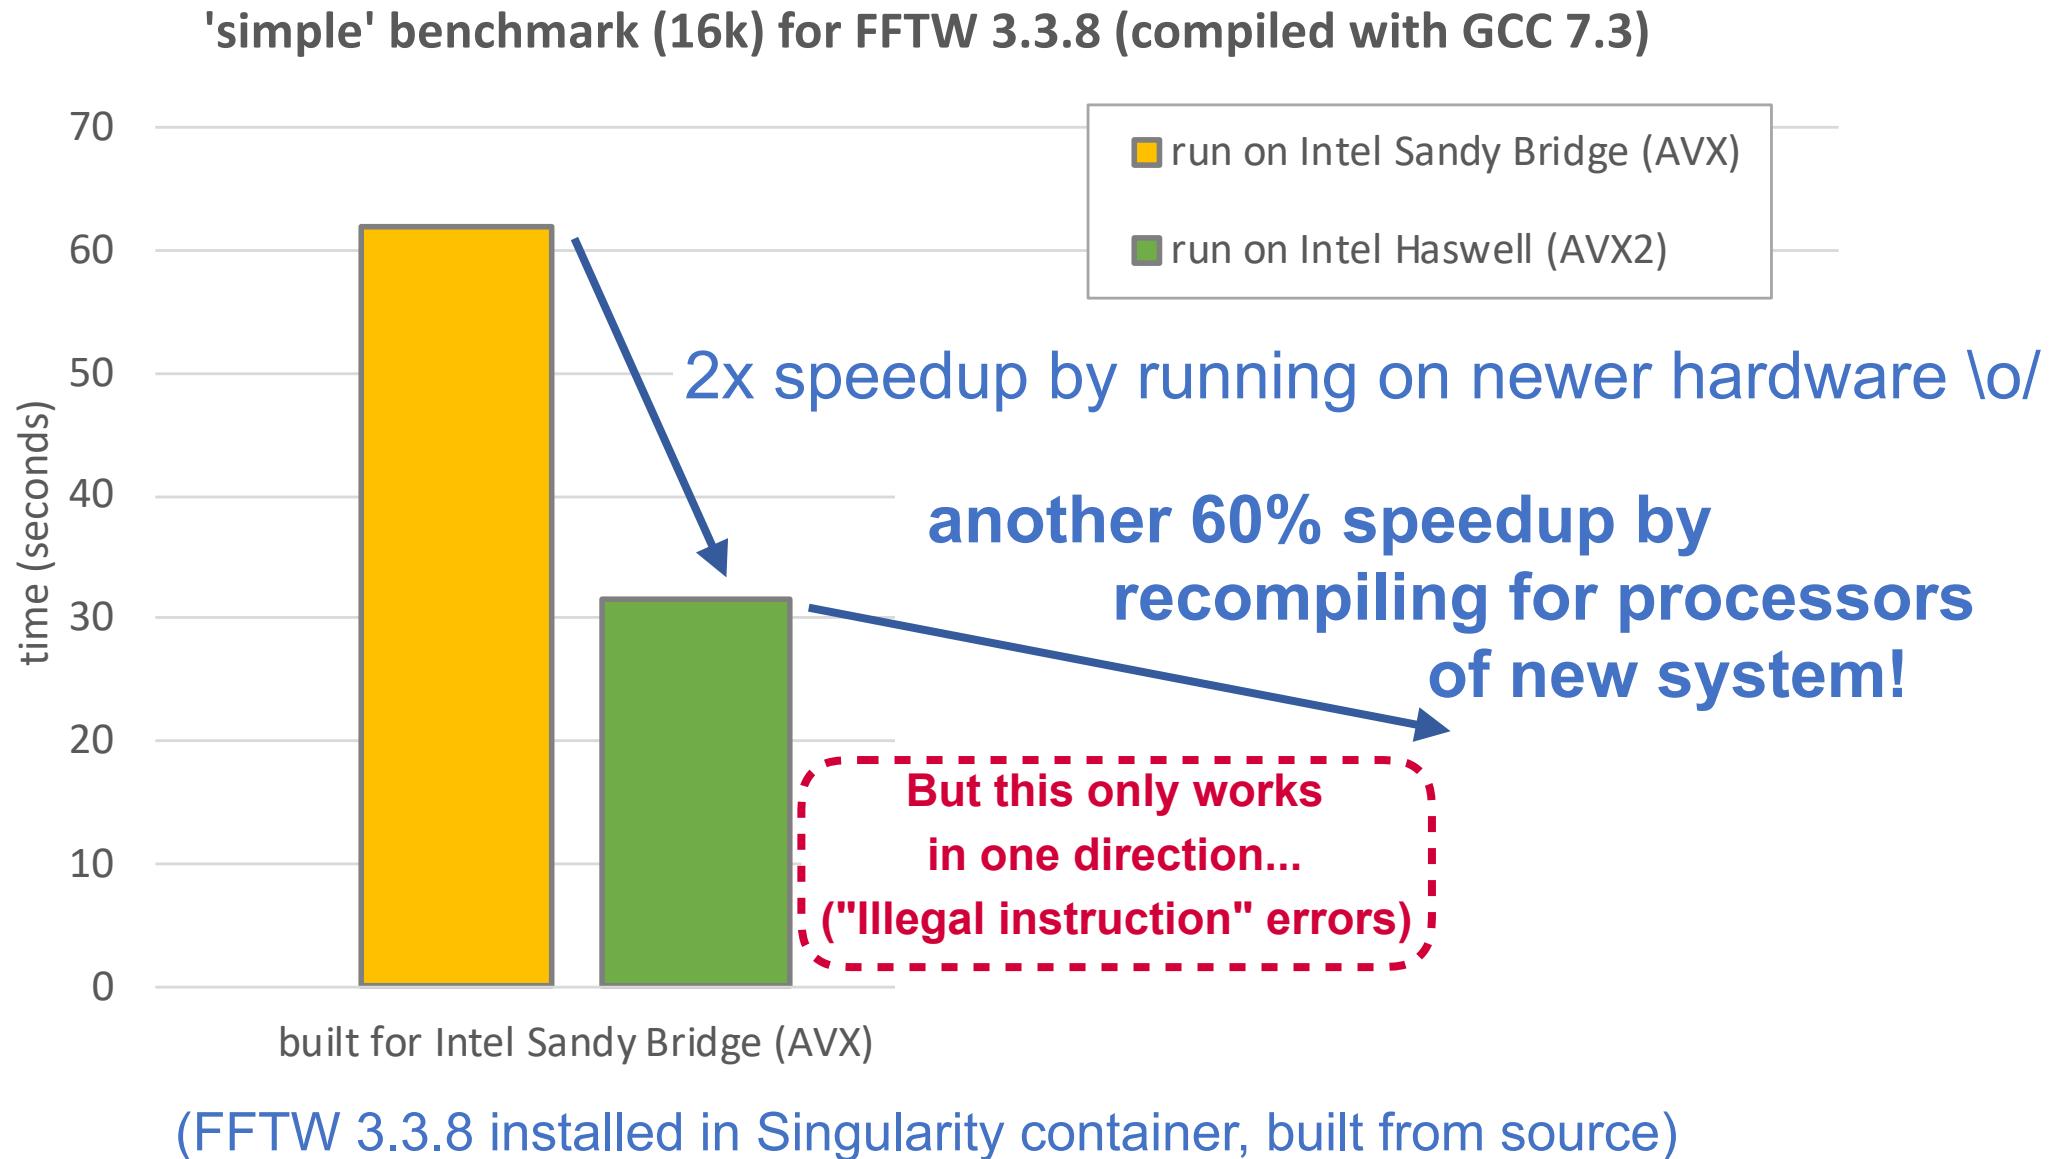
- very popular tool for installing scientific software, but mainly **targeted to individual users**
- **not well suited for centrally managing software** installations on multi-user HPC systems
- installing software with conda *usually* results in running **generically optimised binaries**
- hard to combine with software installed in other ways (e.g. centrally provided modules)

Containers (Docker/Singularity)

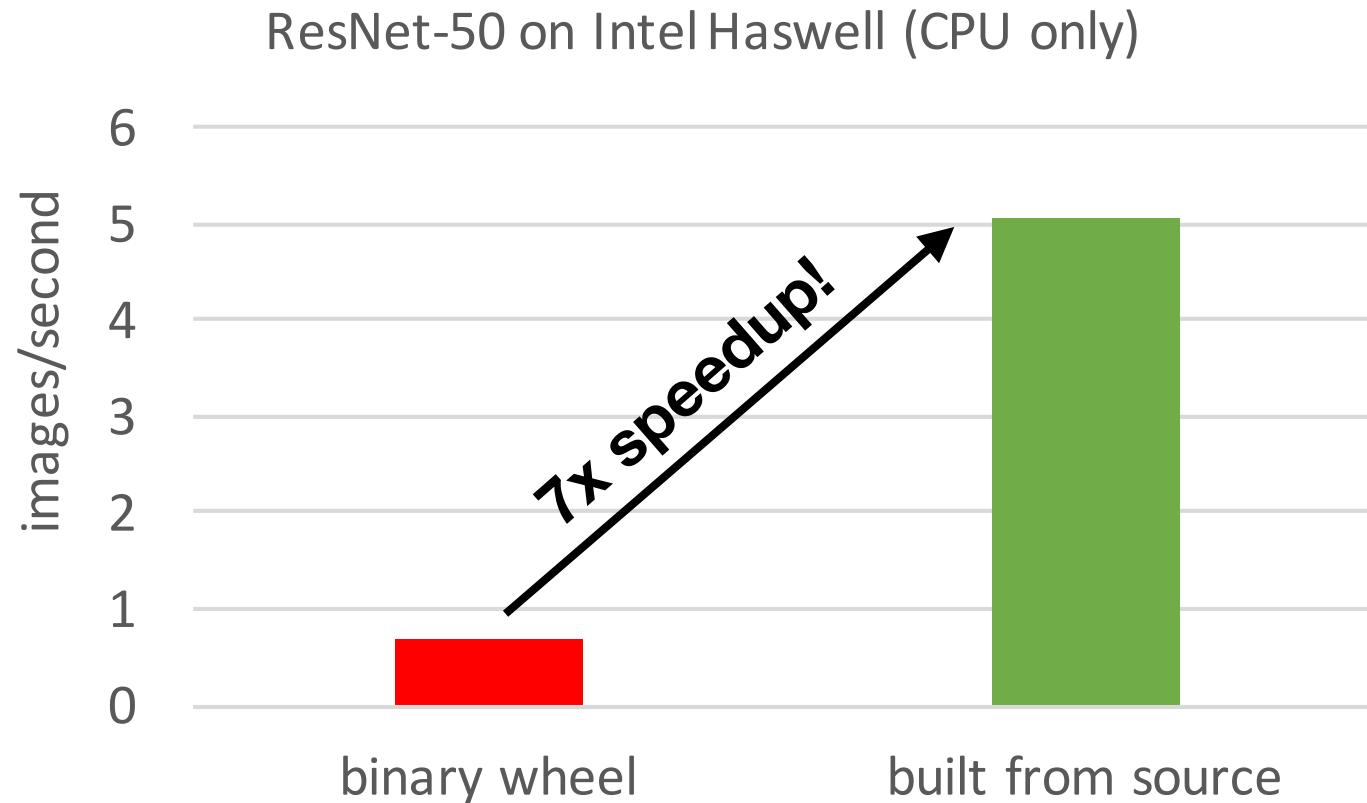


- **lack of attention to optimising for underlying hardware** ("mobility of compute")
- **integration with system resources** (MPI, CUDA, ...) is ***still a problem***
- **someone still needs to be build/maintain/update the container image you need/want !**
- in my opinion, containers are a *symptom*, not a cure...

Performance is key in HPC



"pip install tensorflow" works fine for me...



(disclaimer: old version of TensorFlow (1.4.x?), but point is still valid)

- installing pre-built generically optimised binaries is easy
- ... but **you may be paying a *significant* prize w.r.t. performance**

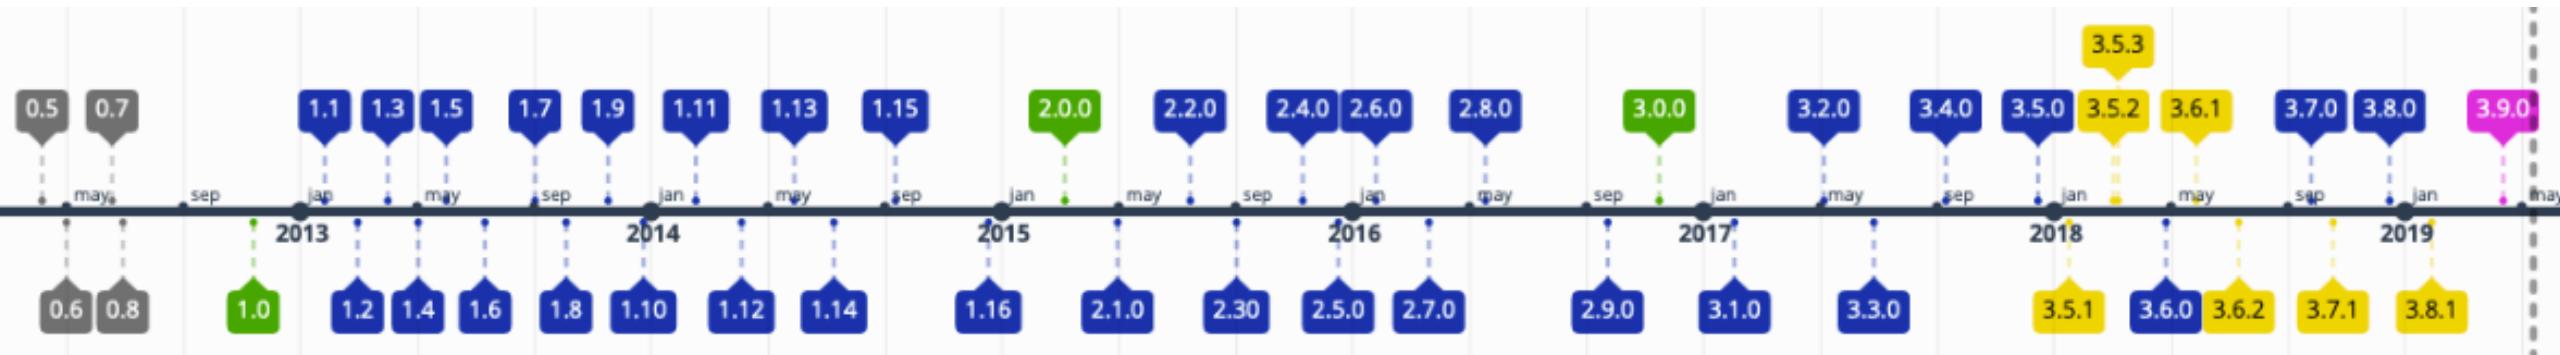


<https://easybuilders.github.io/easybuild> - <https://easybuild.readthedocs.io>

- **framework for installing scientific software** (*built from source when possible*)
- strong focus on Linux & HPC systems (and hence also performance)
- **by default builds/optimises specifically for host architecture**
- implemented in Python 2, lead development by HPC-UGent
- available under GPLv2 license via PyPI, GitHub
- supports different compilers & MPI libraries, x86_64/ARM/POWER, ...
- **good integration with Cray Programming Environment**
- active & helpful worldwide community

Almost 10 years of easybuild

- in-house development at HPC-UGent since summer 2009
- first public release in April 2012 (EasyBuild v0.5)
- **first stable release on November 13th 2012, during SC'12 (EasyBuild v1.0)**
- intention was to get feedback, but gradually a community emerged around it...
- **frequent stable releases** since then (latest: EasyBuild v3.9.0, Apr 12th 2019)
- community-driven development: bug reports, feature requests, contributions

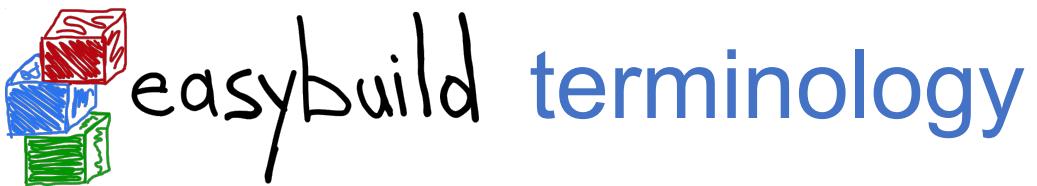


Supported software



http://easybuild.readthedocs.io/en/latest/version-specific/Supported_software.html

- latest EasyBuild (v3.9.0) supports installing **over 1,700 different software packages**
 - including CP2K, NAMD, NWChem, OpenFOAM, TensorFlow, WRF, ...
 - a lot of bioinformatics software is also supported out of the box
 - + >1,000 extensions: Python packages, R libraries, Perl modules, X11 libraries, ...
 - built from source when possible, optimised by host architecture by default
- diverse toolchain support:
 - compilers: GCC, Intel, Clang, PGI, IBM XL, **Cray GNU/Intel/CCE**, CUDA
 - MPI libraries: OpenMPI, Intel MPI, MPICH, MPICH2, MVAPICH2, **Cray MPI**, ...
 - BLAS/LAPACK libraries: Intel MKL, OpenBLAS, ScaLAPACK, BLIS, **Cray LibSci**, ...



http://easybuild.readthedocs.io/en/latest/Concepts_and_Terminology.html

- **framework**
 - core of EasyBuild: Python modules & packages
 - provides supporting functionality for building/installing software, generating modules, ...
- **easyblock**
 - a Python module that serves as a build script, 'plugin' for the EasyBuild framework
 - implements a (generic or software-specific) build/install procedure
- **easyconfig file (*.eb)**: build specification; software name/version, compiler toolchain, etc.
- **(compiler) toolchain**: set of compilers + accompanying libraries (MPI, BLAS/LAPACK, ...)
- **extensions**: additional packages for a particular applications (e.g., Python, R)



easybuild feature highlights (1)

- fully **autonomously** building and installing (scientific) software
 - automatic dependency resolution (via `--robot`)
 - automatic generation of environment module files (Tcl or Lua syntax)
- **no admin privileges required** (only write permission to installation target)
- thorough **logging** of executed build/install procedure
- **archiving** of easyconfigs, patches, easyblocks that were used
- highly **configurable**, via config files/environment/command line
- **dynamically extendable** with additional easyblocks, toolchains, etc.



easybuild feature highlights (2)

- support for **custom module naming schemes** (incl. hierarchical)
- **transparency** via support for 'dry run' installation & trace output
- **comprehensively tested**: lots of unit tests, frequent regression testing, ...
- actively developed, **frequent stable releases**
- **collaboration** between various HPC sites large & small
- integration with Torque/SLURM, FPM, Singularity, Docker, ...
- worldwide **community**

Integration with Cray Programming Environment



<https://easybuild.readthedocs.io/en/latest/Cray-support.html>

- initial (experimental) integration with Cray PE implemented April 2015
- declared stable in March 2016 (EasyBuild v2.7.0)
- custom Cray* toolchains leverage Cray-provided modules:

PrgEnv/*, cray-libsci, cray-mpich, ...

- --optarch controls which craype-* module is used in build environment
- leveraging of Cray-provided "external" modules via provided metadata
- careful handling of Cray-provided modules in build environment & toolchains
- **used by CSCS since 2015 on Piz Daint & other Cray systems**

What easybuild is not

- EasyBuild is **not** YABT (Yet Another Build Tool)

it does *not* replace build tools like `cmake` or `make`; it wraps around them

- it is **not** a replacement for package managers (yum, apt, ...)

it leverages some tools & libraries provided by the OS (glibc, OpenSSL, IB drivers, ...)

- it is **not** a magic solution to all your (software installation) problems...

you will still run into compiler errors (unless somebody has already taken care of it)

What easybuild is

- a **uniform interface** that wraps around software installation procedures
- a huge **time-saver**, by automating tedious/boring/repetitive tasks
- a way to provide a **consistent software stack** to your users
- an **expert system** for software installation on HPC systems
- a **platform for collaboration** with HPC sites worldwide
- a way to **empower users to self-manage their software stack** on HPC systems
- a tool that can be leveraged for **building optimised container images**

Installing easybuild

<https://easybuild.readthedocs.io/en/latest/Installation.html>

EasyBuild is already installed on Piz Daint!

\$ module load daint-mc EasyBuild-custom

<https://user.csics.ch/computing/compilation/easybuild>

- requirements:
 - GNU/Linux system + Python 2.6 or 2.7 (Python 3 will be supported soon)
 - environment modules tool (default configuration requires Lmod)
 - a system C/C++ compiler (only really needed to install toolchains)
- installation procedure (pick one):
 - `pip install easybuild` (or equivalent thereof using other Python installation tool)
 - bootstrap script (recommended)

```
$ python bootstrap_eb.py $EASYBUILD_PREFIX  
$ module use $EASYBUILD_PREFIX/modules/all  
$ module load EasyBuild
```

Updating easybuild

<https://easybuild.readthedocs.io/en/latest/Installation.html#updating-an-existing-easybuild-installation>

- new versions of EasyBuild 3.x are a drop-in replacement (for older EasyBuild 3.x)
 - upcoming EasyBuild 4.0 will have some minor backwards-incompatible changes
- recommend way to pick up new EasyBuild version & start using it:

`eb --install-latest-eb-release`

```
kehoste@daint103:~> module load daint-gpu EasyBuild-custom
kehoste@daint103:~> eb --version
This is EasyBuild 3.8.1 (framework: 3.8.1, easyblocks: 3.8.1) on host daint103.
```

```
kehoste@daint103:~> eb --install-latest-eb-release
kehoste@daint103:~> module load EasyBuild/3.9.0
kehoste@daint103:~> eb --version
This is EasyBuild 3.9.0 (framework: 3.9.0, easyblocks: 3.9.0) on host daint103.
kehoste@daint103:~> which eb
/users/kehoste/easybuild/daint/haswell/software/EasyBuild/3.9.0/bin/eb
```

Configuring easybuild

<http://easybuild.readthedocs.io/en/latest/Configuration.html>

EasyBuild is pre-configured on Piz Daint!

```
$ module load daint-mc EasyBuild-custom
```

```
https://user.cscs.ch/computing/compilation/easybuild
```

- default configuration (ab)uses `$HOME/.local/easybuild`
- configuring can be done via **config files**, **environment**, or **command line interface**
 - all configuration options are supported on all 3 levels
 - CLI overrides environment, which overrides configuration files
- easy to extend EasyBuild:
 - (also) use own easyconfigs repositories via `--robot-paths`
 - add additional easyblocks, toolchains, module naming schemes via `--include-*`
 - http://easybuild.readthedocs.io/en/latest/Including_additional_Python_modules.html

Inspecting the current configuration for easybuild

<http://easybuild.readthedocs.io/en/latest/Configuration.html>

- use 'eb --show-config' to get an overview of the current configuration
- only shows a couple of important settings + anything different from default
- output indicates via which configuration level each setting was specified
- full list of settings for current configuration via 'eb --show-full-config'

```
$ EASYBUILD_PREFIX=/tmp eb --buildpath /dev/shm --show-config
#
# Current EasyBuild configuration
# (C: command line argument, D: default value, E: environment variable, F: configuration file)
#
buildpath      (C) = /dev/shm
installpath    (E) = /tmp
packagepath    (E) = /tmp/packages
prefix         (E) = /tmp
repositorypath (E) = /tmp/ebfiles_repo
robot-paths    (D) = /home/example/easybuild-easyconfigs/easybuild/easyconfigs
sourcepath     (E) = /tmp/sources
```

Default easybuild configuration on Piz Daint

(note: actual output has been trimmed for sake of clarity)

```
$ module load daint-mc EasyBuild-custom

$ eb --show-config
#
# Current EasyBuild configuration
# (C: command line argument, D: default value, E: environment variable, F: configuration file)
#
buildpath          (E) = /run/user/23189/easybuild/build
external-modules-metadata (E) = /apps/common/UES/jenkins/production/easybuild/cray_external_...
hide-deps          (F) = absl, ANTLR, APR, APR-util, Autoconf, Automake, Autotools, ...
hide-toolchains    (F) = CrayCCE, CrayGNU, CrayIntel, CrayPGI, GCCcore, ...
include-easyblocks (E) = /apps/common/UES/jenkins/production/easybuild/easyblocks/*.py
installpath         (E) = /users/kehoste/easybuild/daint/broadwell
module-syntax      (F) = Tcl
modules-tool        (E) = EnvironmentModuleSC
optarch             (E) = broadwell
prefix              (E) = /users/kehoste/easybuild/daint/broadwell
repositorypath     (E) = /users/kehoste/easybuild/daint/broadwell/ebfiles_repo
robot-paths         (E) = /apps/common/UES/jenkins/production/easybuild/easyconfigs/, ...
sourcepath          (E) = /users/kehoste/sources
tmpdir              (E) = /run/user/23189/easybuild/tmp
```

Default easybuild configuration on Piz Daint

(note: actual output has been trimmed for sake of clarity)

```
$ module load daint-mc EasyBuild-custom
```

```
$ eb --show-config
```

User-specific '/run' directory is used for build and temporary directories.

```
#  
# Current EasyBuild configuration  
# (C: command line argument, D: default value, E: environment variable, F: configuration file)  
  
buildpath (E) = /run/user/23189/easybuild/build  
external-modules-metadata (E) = /apps/common/UES/jenkins/production/easybuild/cray_external_...  
hide-deps (F) = absl, ANTLR, APR, APR-util, Autoconf, Automake, Autotools, ...  
hide-toolchains (F) = CrayCCE, CrayGNU, CrayIntel, CrayPGI, GCCcore, ...  
include-easyblocks (E) = /apps/common/UES/jenkins/production/easybuild/easyblocks/*.py  
installpath (E) = /users/kehoste/easybuild/daint/broadwell  
module-syntax (F) = Tcl  
modules-tool (E) = EnvironmentModuleSC  
optarch (E) = broadwell  
prefix (E) = /users/kehoste/easybuild/daint/broadwell  
repositorypath (E) = /users/kehoste/easybuild/daint/broadwell/ebfiles_repo  
robot-paths (E) = /apps/common/UES/jenkins/production/easybuild/easyconfigs/, ...  
sourcepath (E) = /users/kehoste/sources  
tmpdir (E) = /run/user/23189/easybuild/tmp
```

Default easybuild configuration on Piz Daint

(note: actual output has been trimmed for sake of clarity)

```
$ module load daint-mc EasyBuild-custom
$ eb --show-config
#
# Current EasyBuild configuration
# (C: command line argument, D: default value, E: environment variable, F: configuration file)
#
buildpath          (E) = /run/user/23189/easybuild/build
external-modules-metadata (E) = /apps/common/UES/jenkins/production/easybuild/cray_external_...
hide-deps          (F) = absl, ANTLR, APR, APR-util, Autoconf, Automake, Autotools, ...
hide-toolchains    (F) = CrayCCE, CrayGNU, CrayIntel, CrayPGI, GCCcore, ...
include-easyblocks (E) = /apps/common/UES/jenkins/production/easybuild/easyblocks/*.py
installpath         (E) = /users/kehoste/easybuild/daint/broadwell
module-syntax      (F) = Tcl
modules-tool        (E) = EnvironmentModuleSC
optarch             (E) = broadwell
prefix              (E) = /users/kehoste/easybuild/daint/broadwell
repositorypath     (E) = /users/kehoste/easybuild/daint/broadwell/ebfiles_repo
robot-paths         (E) = /apps/common/UES/jenkins/production/easybuild/easyconfigs/, ...
sourcepath           (E) = /users/kehoste/sources
tmpdir              (E) = /run/user/23189/easybuild/tmp
```

Default easybuild configuration on Piz Daint

(note: actual output has been trimmed for sake of clarity)

```
$ module load daint-mc EasyBuild-custom
$ eb --show-config
#
# Current EasyBuild configuration
# (C: command line argument, D: default value, E: environment variable, F: configuration file)
#
buildpath          (E) = /run/user/23189/easybuild/build
external-modules-metadata (E) = /apps/common/UES/jenkins/production/easybuild/cray_external_...
hide-deps          (F) = absl, ANTLR, APR, APR-util, Autoconf, Automake, Autotools, ...
hide-toolchains     (F) = CrayCCE, CrayGNU, CrayIntel, CrayPGI, GCCcore, ...
include-easyblocks (E) = /apps/common/UES/jenkins/production/easybuild/easyblocks/*.py
installpath         (E) = /users/kehoste/easybuild/daint/broadwell
module-syntax       (F) = Tcl
modules-tool        (E) = EnvironmentModulesC
optarch             (E) = broadwell
prefix              (E) = /users/kehoste/easybuild/daint/broadwell
repositorypath      (E) = /users/kehoste/easybuild/daint/broadwell/ebfiles_repo
robot-paths         (E) = /apps/common/UES/jenkins/production/easybuild/easyconfigs/, ...
sourcepath           (E) = /users/kehoste/sources
tmpdir              (E) = /run/user/23189/easybuild/tmp
```

Custom configuration for Piz Daint (1/2):

- metadata for Cray-provided external modules
- Tcl-based environment modules tool (Lmod is default)
- target optimisation architecture ('broadwell' or 'haswell')

Default easybuild configuration on Piz Daint

(note: actual output has been trimmed for sake of clarity)

```
$ module load daint-mc EasyBuild-custom
$ eb --show-config
#
# Current EasyBuild configuration
# (C: command line argument, D: default value, E: environment variable, F: configuration file)
#
buildpath          (E) = /run/user/23189/easybuild/build
external-modules-metadata (E) = /apps/common/UES/jenkins/production/easybuild/cray_external_...
hide-deps          (F) = absl, ANTLR, APR, APR-util, Autoconf, Automake, Autotools, ...
hide-toolchains    (F) = CrayCCE, CrayGNU, CrayIntel, CrayPGI, GCCcore, ...
include-easyblocks (E) = /apps/common/UES/jenkins/production/easybuild/easyblocks/*.py
installpath         (E) = /users/kehoste/easybuild/daint/broadwell
module-syntax      (F) = Tcl
modules-tool        (E) = EnvironmentModuleSC
optarch             (E) = broadwell
prefix              (E) = /users/kehoste/easybuild/daint/broadwell
repositorypath     (E) = /users/kehoste/easybuild/daint/broadwell/ebfiles_repo
robot-paths         (E) = /apps/common/UES/jenkins/production/easybuild/easyconfigs/, ...
sourcepath          (E) = /users/kehoste/sources
tmpdir              (E) = /run/user/23189/easybuild/tmp
```

Custom configuration for Piz Daint (2/2):

- hiding dependencies & toolchain (in module avail)
- customised easyblocks maintained by CSCS
- additional easyconfig files (using Cray* toolchains)

Basic usage



http://easybuild.readthedocs.io/en/latest/Using_the_EasyBuild_command_line.html

http://easybuild.readthedocs.io/en/latest/Typical_workflow_example_with_WRF.html

- specify software name/version and toolchain to 'eb' command
- usually done via easyconfig filename(s):

```
eb GCC-7.3.0-2.30.eb Clang-7.01-7.3.0-2.30.eb
```

- check whether required toolchain & dependencies are available using --dry-run/-D:

```
eb Python-3.7.0-foss-2018b.eb -D
```

- enable dependency resolution via --robot/-r:

```
eb TensorFlow-1.13.1-foss-2018b-Python-3.6.6.eb -r
```

Installing TensorFlow from source with **one command**...

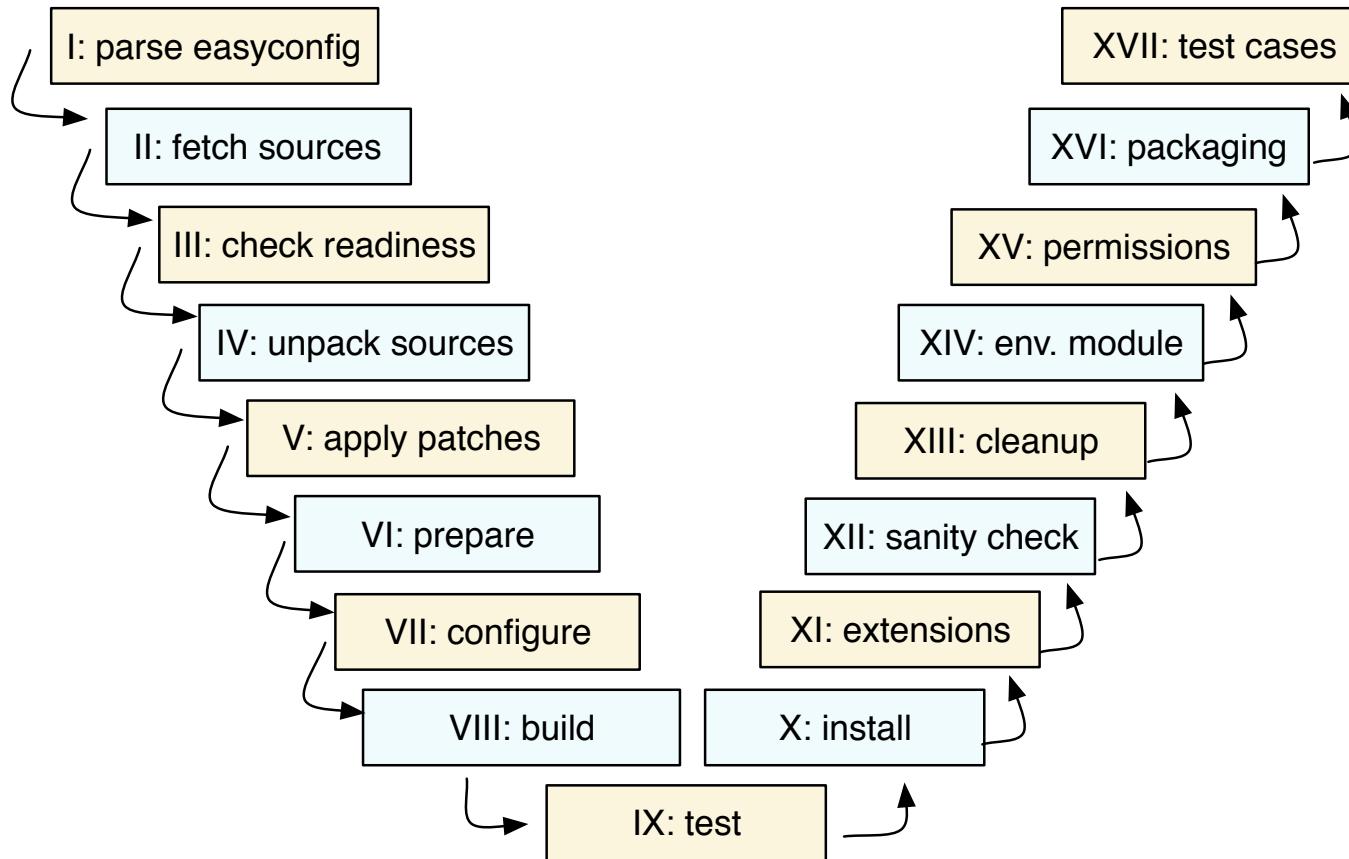


```
$ eb TensorFlow-1.13.1-foss-2018b-Python-3.6.6.eb
== temporary log file in case of crash /tmp/eb-GyvPHx/easybuild-U1TkEI.log
== processing EasyBuild easyconfig TensorFlow-1.13.1-foss-2018b-Python-3.6.6.eb
== building and installing TensorFlow/1.13.1-foss-2018b-Python-3.6.6...
== fetching files...
== creating build dir, resetting environment...
== unpacking...
== patching...
== preparing...
== configuring...
== building...
== testing...
== installing...
== taking care of extensions...
== postprocessing...
== sanity checking...
== cleaning up...
== creating module...
== permissions...
== packaging...
== COMPLETED: Installation ended successfully
== Results of the build can be found in the log file /opt/easybuild/software/Tensor...
== Build succeeded for 1 out of 1
== Temporary log file(s) /tmp/eb-GyvPHx/easybuild-U1TkEI.log* have been removed.
== Temporary directory /tmp/eb-GyvPHx has been removed.
```

Step-wise installation procedure



EasyBuild performs a step-wise installation procedure for each software:



- download sources (best effort)
- set up build directory & environment
 - unpack sources (& apply patches)
 - load modules for toolchain & deps
 - define toolchain-related env vars (\$CC, \$CFLAGS, ...)
- configure, build, (test), install, (extensions)
- perform simple sanity check on installation
- generate environment module file

each step can be customised via easyconfig parameters or an easyblock

Transparency of performed install procedure (1)



https://easybuild.readthedocs.io/en/latest/Tracing_progress.html

eb --trace shows more details while EasyBuild is performing installation(s)

- exact location of build and install directories
- list of source files & patches
- modules being loaded
- commands being executed
- pointer to temporary log file with output of commands
- timing information for each command (start time + how long it took)
- results of sanity check

Example output of eb --trace



```
$ eb TensorFlow-1.13.1-foss-2018b-Python-3.6.6.eb --trace
...
== preparing...
>> loading toolchain module: foss/2018b
>> loading modules for build dependencies:
>> * Bazel/0.20.0-GCCcore-7.3.0
>> * protobuf/3.6.1-GCCcore-7.3.0
>> loading modules for (runtime) dependencies:
>> * Python/3.6.6-foss-2018b
>> * wheel/0.31.1-foss-2018b-Python-3.6.6
>> * h5py/2.8.0-foss-2018b-Python-3.6.6
>> defining build environment for foss/2018b toolchain
...
== installing extension TensorFlow 1.13.1 (13/13)...
...
>> running command:
    [started at: 2019-05-12 14:12:38]
    [output logged in /tmp/eb-pRHwkc/easybuild-run_cmd-SOINRV.log]
        bazel ... --jobs=6 --config=mkl //tensorflow/tools/pip_package:build_pip_package
>> command completed: exit 0, ran in 00h41m22s
...
```

Transparency of performed install procedure (2)



https://easybuild.readthedocs.io/en/latest/Extended_dry_run.html

- **eb --extended-dry-run** (or '**eb -x**') reveals planned installation procedure
- runs in a matter of seconds
- shows commands that will be executed, build environment, generated module file, ...
- any errors that occur in used easyblock are ignored (but clearly reported)
- not 100% accurate since easyblock may require certain files to be present, etc.
- very useful when debugging easyblocks, instant feedback as a first pass
- implementation motivated by requests from the community
- helps to avoid impression that EasyBuild is a magic black box for installing software

Example output of --extended-dry-run (1/3)



```
$ eb WRF-3.8.0-intel-2016b-dmpar.eb -x
```

```
== temporary log file in case of crash /tmp/eb-Dh1wOp/easybuild-0bu9u9.log
```

```
== processing EasyBuild easyconfig /home/example/eb/easybuild-easyconfigs/easybuild/
easyconfigs/w/WRF/WRF-3.8.0-intel-2016b-dmpar.eb
```

```
...
```

```
*** DRY RUN using 'EB_WRF' easyblock (easybuild.easyblocks.wrf @ /home/example/eb/easybuild-
easyblocks/easybuild/easyblocks/w/wrf.py) ***
```

```
== building and installing WRF/3.8.0-intel-2016b-dmpar...
fetching files... [DRY RUN]
```

```
[fetch_step method]
```

```
Available download URLs for sources/patches:
```

- * [http://www2.mmm.ucar.edu/wrf/src//\\$source](http://www2.mmm.ucar.edu/wrf/src//$source)
- * [http://www.mmm.ucar.edu/wrf/src//\\$source](http://www.mmm.ucar.edu/wrf/src//$source)

```
List of sources:
```

- * WRFV3.8.0.TAR.gz will be downloaded to /home/example/eb/sources/w/WRF/WRFV3.8.0.TAR.gz

Example output of --extended-dry-run (2/3)



```
$ eb WRF-3.8.0-intel-2016b-dmpar.eb -x
```

```
...
```

```
building... [ DRY RUN ]
```

```
[build_step method]
```

```
running command "tcsh ./compile -j 4 wrf"  
(in /home/example/eb/software/WRF/3.8.0-intel-2016b-dmpar/WRF-3.8.0)
```

```
running command "tcsh ./compile -j 4 em_real"
```

```
(in /home/example/eb/software/WRF/3.8.0-intel-2016b-dmpar/WRF-3.8.0)
```

```
running command "tcsh ./compile -j 4 em_b_wave"
```

```
(in /home/example/eb/software/WRF/3.8.0-intel-2016b-dmpar/WRF-3.8.0)
```

```
...
```

```
[sanity_check_step method]
```

```
Sanity check paths - file ['files']
```

```
* WRFV3/main/libwrflib.a
```

```
* WRFV3/main/real.exe
```

```
* WRFV3/main/wrf.exe
```

```
Sanity check paths - (non-empty) directory ['dirs']
```

```
* WRFV3/main
```

```
* WRFV3/run
```

```
Sanity check commands
```

```
(none)
```

Example output of --extended-dry-run (3/3)



```
$ eb WRF-3.8.0-intel-2016b-dmpar.eb -x
```

```
...
```

```
[make_module_step method]
```

```
Generating module file /home/example/eb/modules/all/WRF/3.8.0-intel-2016b-dmpar,  
with contents:
```

```
#%Module  
proc ModulesHelp { } {  
    puts stderr { The Weather Research and Forecasting (WRF) Model }  
}  
module-whatis {Description: WRF - Homepage: http://www.wrf-model.org}  
  
set root /home/example/eb/software/WRF/3.8.0-intel-2016b-dmpar  
  
conflict WRF  
  
if { ![ is-loaded intel/2016b ] } {  
    module load intel/2016b  
}  
if { ![ is-loaded JasPer/1.900.1-intel-2016b ] } {  
    module load JasPer/1.900.1-intel-2016b  
}
```

Log files



<http://easybuild.readthedocs.io/en/latest/Logfiles.html>

- **EasyBuild *thoroughly* logs the executed installation procedure**
 - active EasyBuild configuration
 - easyconfig file that was used
 - modules that were loaded + resulting changes to environment
 - defined environment variables
 - output + exit code of executed commands
 - informative log messages produced by easyblock
- log file is copied to software installation directory for future reference
- can be used to debug build problems or see how installation was performed exactly

Log files: example



```
== 2019-05-13 13:34:31,906 main.EB_HPL INFO This is EasyBuild 3.9.0 (framework:  
3.9.0, easyblocks: 3.9.0) on host example.  
...  
== 2019-05-13 13:34:35,503 main.EB_HPL INFO configuring...  
== 2019-05-13 13:34:48,817 main.EB_HPL INFO Starting configure step  
...  
== 2019-05-13 13:34:48,823 main.EB_HPL INFO Running method configure_step part of  
step configure  
...  
== 2019-05-13 13:34:48,823 main.run DEBUG run_cmd: running cmd /bin/bash  
make_generic (in /tmp/easybuild_build/HPL/2.3/foss-2019a/hpl-2.3/setup)  
== 2019-05-13 13:34:48,823 main.run DEBUG run_cmd: Command output will be logged  
to /tmp/easybuild-W85p4r/easybuild-run_cmd-XoJwMY.log  
== 2019-05-13 13:34:48,849 main.run INFO cmd "/bin/bash make_generic" exited with  
exitcode 0 and output:  
...
```

Easyconfig files as build specifications



http://easybuild.readthedocs.io/en/latest/Writing_easyconfig_files.html

- for each software installation, there is a corresponding easyconfig file
- **simple text files defining a set of easyconfig parameters** (in Python syntax)
- some are mandatory: software name/version, toolchain, metadata (homepage, descr.)
- other commonly used parameters:
 - easyblock to use
 - list of sources & patches
 - list of (build) dependencies
 - options for configure/build/install commands
 - files/directories that should be present, trivial commands that should work (sanity check)

Overview of supported easyconfig parameters:

\$ eb -a

Include easyconfig parameters specific to a particular a particular easyblock:

\$ eb -a -e PythonPackage

Example easyconfig file



no easyblock specified, which implies using a software-specific easyblock (EB_WRF)

software name and version

build variant (specific to WRF)
('dmpar': distributed, MPI)

software metadata

toolchain name & version

sources & patches

list of (build) dependencies
note: all versions are *fixed!*

```
name = 'WRF'
version = '3.8.0'
buildtype = 'dmpar'    # custom parameter for WRF
versionsuffix = '-' + buildtype    # part of module name

homepage = 'http://www.wrf-model.org'
description = "Weather Research and Forecasting (WRF) Model"

toolchain = {'name': 'intel', 'version': '2016b'}

source_urls = ['http://www.mmm.ucar.edu/wrf/src/']
sources = ['%(name)sV%(version_major_minor)s.TAR.gz']
patches = ['WRF-%(version)s_known_problems.patch']

builddependencies = [('tcsh', '6.20.00')]
dependencies = [
    ('JasPer', '2.0.10'),
    ('netCDF', '4.4.1'),
    ('netCDF-Fortran', '4.4.4'),
]
```

Adding support for additional software



- for each software installation, an **easyconfig file** is required
 - see https://easybuild.readthedocs.io/en/latest/Writing_easyconfig_files.html
 - existing easyconfig files can serve as examples
 - tweaked easyconfig files can be generated via `eb --try-*`
- for 'standard' installation procedures, a **generic easyblock** can be used
 - installation can be controlled where needed via easyconfig parameters
- for custom installation procedures, a **software-specific easyblock** is required
 - see <https://easybuild.readthedocs.io/en/latest/Implementing-easyblocks.html>

Tweaking existing easyconfigs via --try-*



- simple common use cases include:
 - updating to a more recent software version
 - installing with a different compiler toolchain
- dedicated CLI options: `--try-software-version`, `--try-toolchain`, ...
- success is not guaranteed... (hence the 'try')

```
kehoste@daint104:~> module load daint-mc EasyBuild-custom/cscs
kehoste@daint104:~> module avail CMake
----- /apps/daint/UES/jenkins/6.0.UP07/mc/easybuild/modules/all -----
CMake/3.12.0
kehoste@daint104:~> eb CMake-3.12.0.eb --try-software-version 3.14.3
...
kehoste@daint104:~> module avail CMake
----- /users/kehoste/easybuild/daint/broadwell/modules/all -----
CMake/3.14.3
----- /apps/daint/UES/jenkins/6.0.UP07/mc/easybuild/modules/all -----
CMake/3.12.0
```

Writing easyconfig files



https://easybuild.readthedocs.io/en/latest/Writing_easyconfig_files.html

- set of values for easyconfig parameters, written in Python syntax
(strictly speaking order doesn't matter, but we maintain a preferred order & style)
- strictly required:
`name, version, toolchain, homepage, description`
- commonly used:
`easyblock, versionsuffix, source_urls, sources,`
`(build)dependencies, (pre)configopts, (pre)buil opts,`
`(pre)installopts, exts_list, sanity_check_paths, moduleclass`

Common generic easyblocks



http://easybuild.readthedocs.io/en/latest/version-specific/generic_easyblocks.html

- **ConfigureMake**
standard '`./configure`' - '`make`' - '`make install`' installation procedure
- **CMakeMake**
same as ConfigureMake, but using '`cmake`' for configuring
- **PythonPackage**
installing an individual Python package ('`python setup.py install`', '`pip install`', ...)
- **PythonBundle**
installing a bundle of Python packages
- **MakeCp**
no (standard) configuration step, build with '`make`', install by copying binaries/libraries
- **Tarball**: just unpack sources and copy everything to installation directory
- **Binary**: run binary installer (specified via '`install_cmd`' easyconfig parameter)

Community (common) toolchains

<http://easybuild.readthedocs.io/en/latest/Common-toolchains.html>

- **intel and foss¹ toolchains** are most commonly used in EasyBuild community
- helps to focus efforts of HPC sites using one or both of these toolchains
- updated twice a year, clear versioning scheme: <year>{a,b} (2017b, 2018a, ...)
- latest version:
 - **foss/2019a**
binutils 2.31.1, GCC 8.2, OpenMPI 3.1.3,
OpenBLAS 0.3.5, FFTW 3.3.8
 - **intel/2019a**
binutils 2.31.1 + GCC 8.2 as base
Intel compilers 2019.1.144, Intel MPI 2018.4.274, Intel MKL 2019.1.144

On Piz Daint, Cray* toolchains are typically used:

CrayCCE/18.08
CrayGNU/18.08
CrayIntel/18.08
CrayPGI/18.08

Easyconfig files vs easyblocks



<http://easybuild.readthedocs.io/en/latest/Implementing-easyblocks.html>

- thin line between using custom easyblock and 'fat' easyconfig with generic easyblock
- custom easyblocks are "do once and forget", **central solution to build peculiarities**
- reasons to consider implementing a software-specific easyblock include:
 - 'critical' values for easyconfig parameters required to make installation succeed
 - toolchain-specific aspects of the build and installation procedure (e.g., configure options)
 - interactive commands that need to be run
 - custom (configure) options for dependencies
 - having to create or adjust specific (configuration) files
 - 'hackish' usage of a generic easyblock

Demo: an easyconfig for Hello built with CMake



```
kehoste@daint104:~> module load daint-mc EasyBuild-custom
kehoste@daint104:~> ls
hello-cmake-1.0.tar.gz
kehoste@daint104:~> vim Hello.eb
kehoste@daint104:~> eb Hello.eb --trace
...
== configuring...
== building...
== testing...
== installing...
...
== sanity checking...
>> file 'bin/hello' found: OK
>> running command 'hello': OK
== cleaning up...
...
== COMPLETED: Installation ended successfully
```

```
kehoste@daint104:~> module load Hello/1.0-CrayGNU-18.08
kehoste@daint104:~> hello
Hello World!
```

Demo: an easyconfig for Hello built with CMake



```
kehoste@daint104:~> cat Hello.eb
easyblock = 'CMakeMake'

name = 'Hello'
version = '1.0'

homepage = 'https://example.com'
description = "Hello world!"

toolchain = { 'name': 'CrayGNU', 'version': '18.08' }

sources = ['hello-cmake-%(version)s.tar.gz']

builddependencies = [ ('CMake', '3.14.3', '', True) ]

sanity_check_paths = {
    'files': ['bin/hello'],
    'dirs': [],
}
sanity_check_commands = ["hello"]

moduleclass = 'tools'
```

Contributing to EasyBuild



<http://easybuild.readthedocs.io/en/latest/Contributing.html>

EasyBuild has improved significantly thanks to the community.

You too can contribute back, by:

- sending feedback
- reporting bugs
- joining the discussion (mailing lists, IRC/Slack, EasyBuild conf calls)
- sharing suggestions/ideas for enhancements & additional features
- contributing easyconfigs, enhancing easyblocks, adding support for new software...
- extending & enhancing documentation

How to contribute easyconfig files



<https://easybuild.readthedocs.io/en/latest/Contributing.html>

- open pull request (PR) to develop branch of easybuild-easyconfigs repository
- keep PRs limited in scope (rule of thumb: max. 10 easyconfigs touched per PR)
- if possible, submit test report to show the easyconfig(s) works for you
- take into account failing unit tests (if any) and comments by reviewer (if any)
- be patient...
- enjoy eternal glory once your pull request has been merged!

GitHub integration in easybuild

https://easybuild.readthedocs.io/en/latest/Integration_with_GitHub.html

- goals:
 - automate contribution workflow
 - avoid direct interaction with git and
- initial implementation added support for:
 - opening new PRs & updating existing PRs
 - downloading & using easyconfigs from a PR
 - uploading test reports to a PR
- limited to easyconfig files, for now...
 - but implemented such that it is easy to extend to easyblocks, and even the EasyBuild framework itself



`eb --new-pr`

`eb --update-pr`

`eb --from-pr`

`eb --upload-test-report`

Opening a pull request in 1, 2, 3



```
$ mv sklearn.eb scikit-learn-0.19.1-intel-2017b-Python-3.6.3.eb  
$ mv scikit-learn*.eb easybuild/easyconfigs/s/scikit-learn  
$ git checkout develop && git pull upstream develop  
$ git checkout -b scikit_learn_0191_intel_2017b  
$ git add easybuild/easyconfigs/s/scikit-learn  
$ git commit -m "{data}[intel/2017b] scikit-learn v0.19.1"  
$ git push origin scikit_learn_0191_intel_2017b
```

+ log into GitHub to actually open the pull request (clickety, clickety...)

one single eb command
no git commands
no GitHub interaction

metadata is automatically derived from easyconfig
saves a lot of time!



eb --new-pr sklearn.eb

The easybuild community



group picture
4th EasyBuild User Meeting
@ UC Louvain (Feb 2019)



- EasyBuild community has been growing rapidly the last couple of years
- used by **hundreds of HPC sites and companies worldwide,**
incl. JSC, CSCS, SURFsara, Pfizer, ...
- very welcoming & supportive to newcomers

mailing list: <https://lists.ugent.be/wws/info/easybuild>
Slack channel: easybuild.slack.com
IRC: #easybuild on FreeNode



easybuild maintainers



Damian Alvarez - @damianam
(Jülich Supercomputing Centre)



Miguel Dias Costa - @migueldiascosta
(National University of Singapore)



Pablo Escobar - @pescobar
(sciCORE, University of Basel)



Kenneth Hoste - @boegel
(HPC-UGent)



Adam Huffman - @verdurin
(Big Data Institute, University of Oxford)



Samuel Moors - @smoors
(Free University of Brussels)



Alan O'Cais - @ocaisa
(Jülich Supercomputing Centre)



Bart Oldeman - @bartoldeman
(ComputeCanada)



Ward Poelmans - @wpoely86
(Free University of Brussels)



Åke Sandgren - @akesandgren
(Umeå University, Sweden)

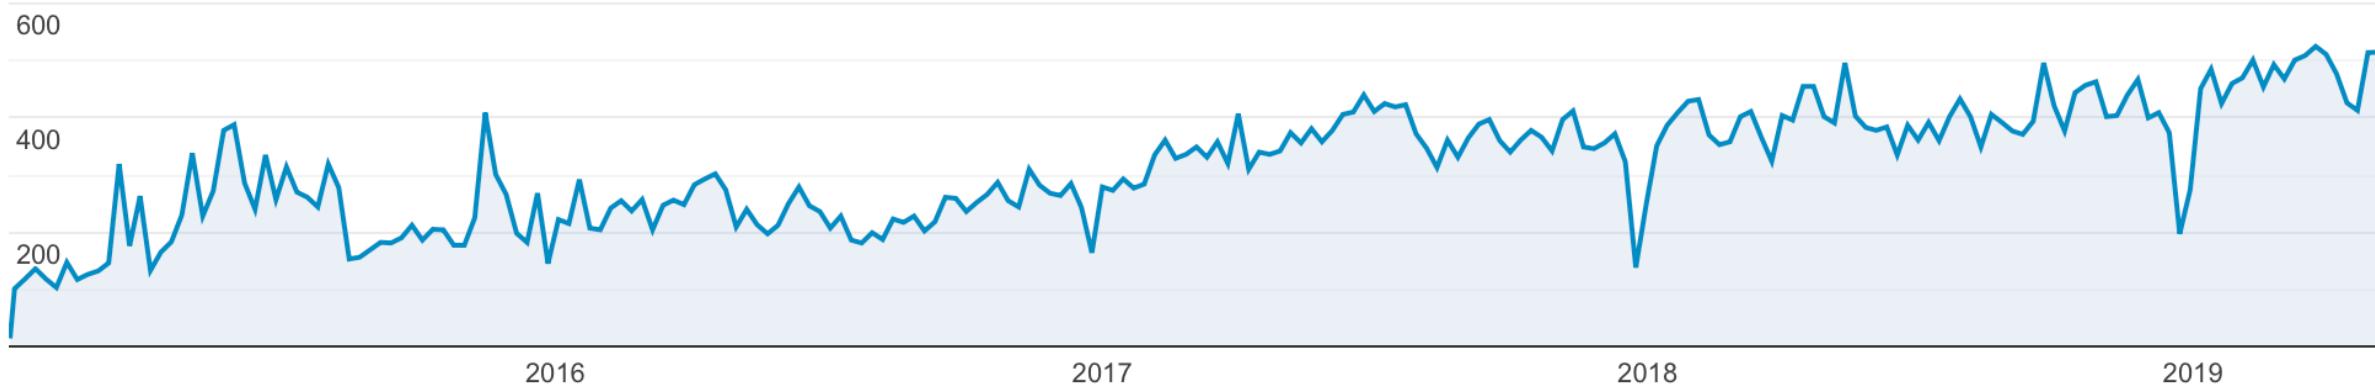


Davide Vanzo - @vanzod
(Vanderbilt University)



Mikael Öhman - @micketeer
(Chalmers University of Technology)

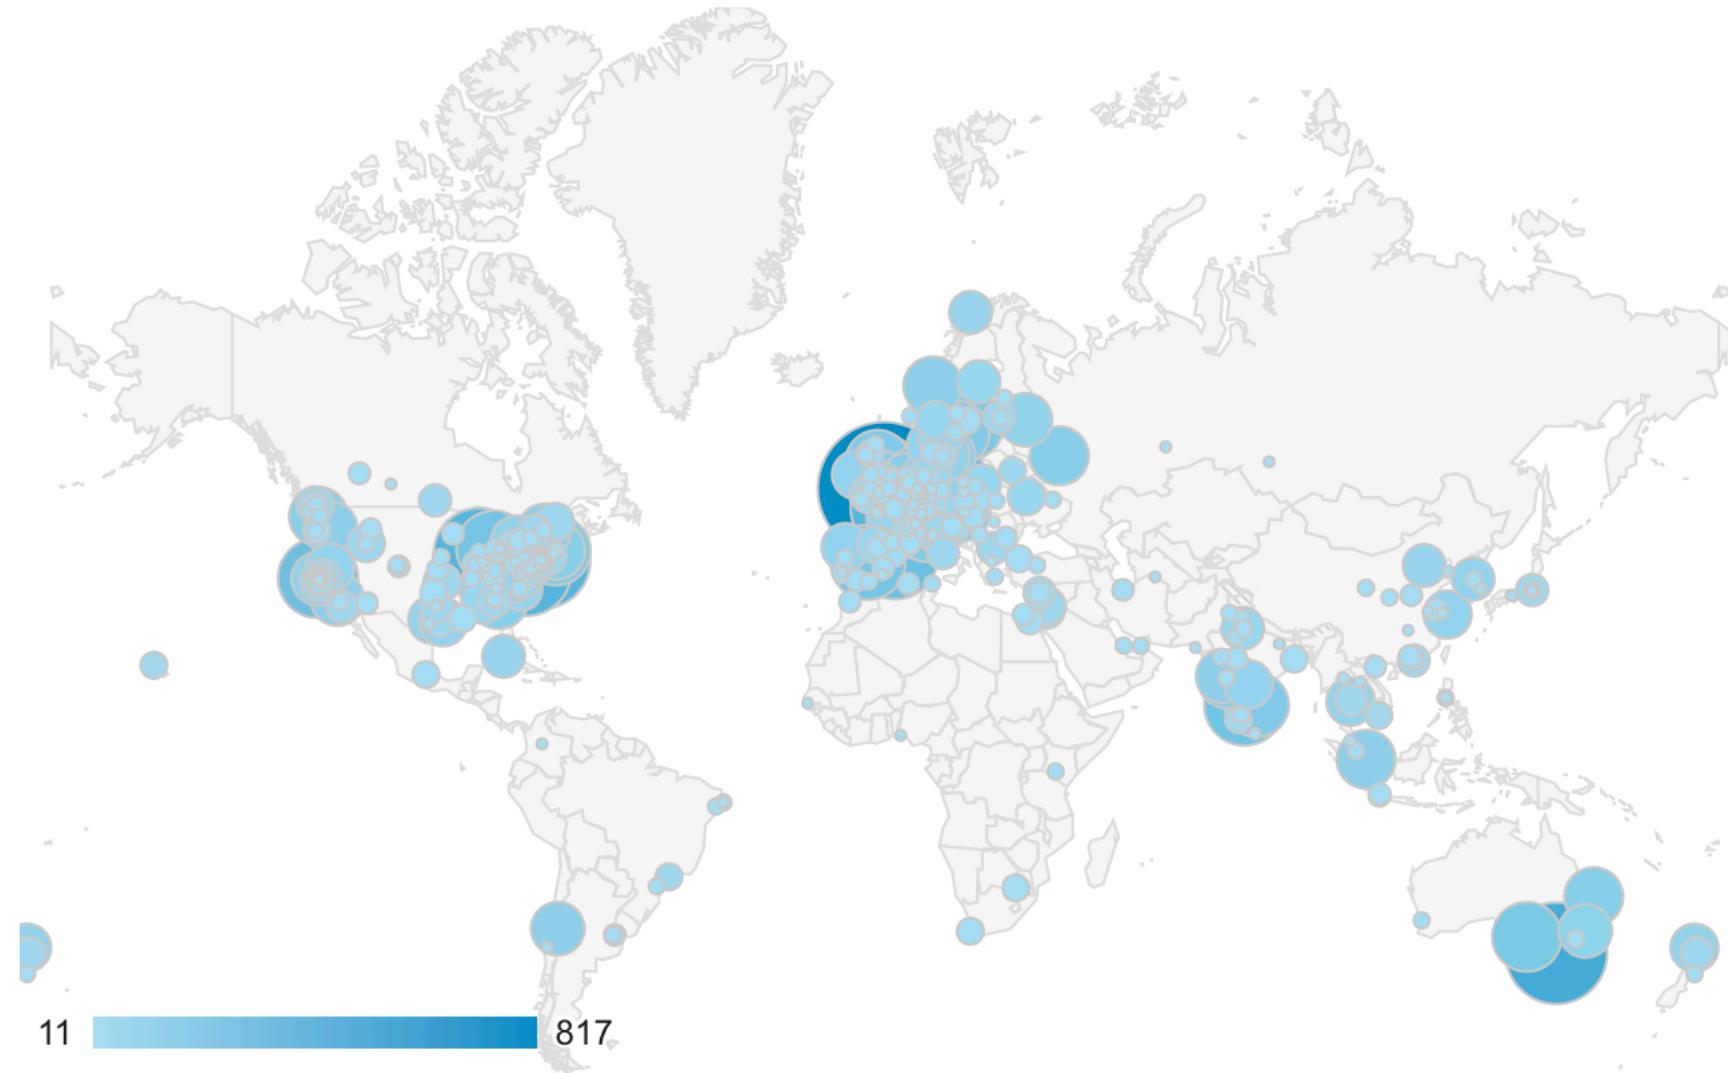
EasyBuild documentation visits



EasyBuild documentation at <https://easybuild.readthedocs.io> hitting ~500 weekly visitors

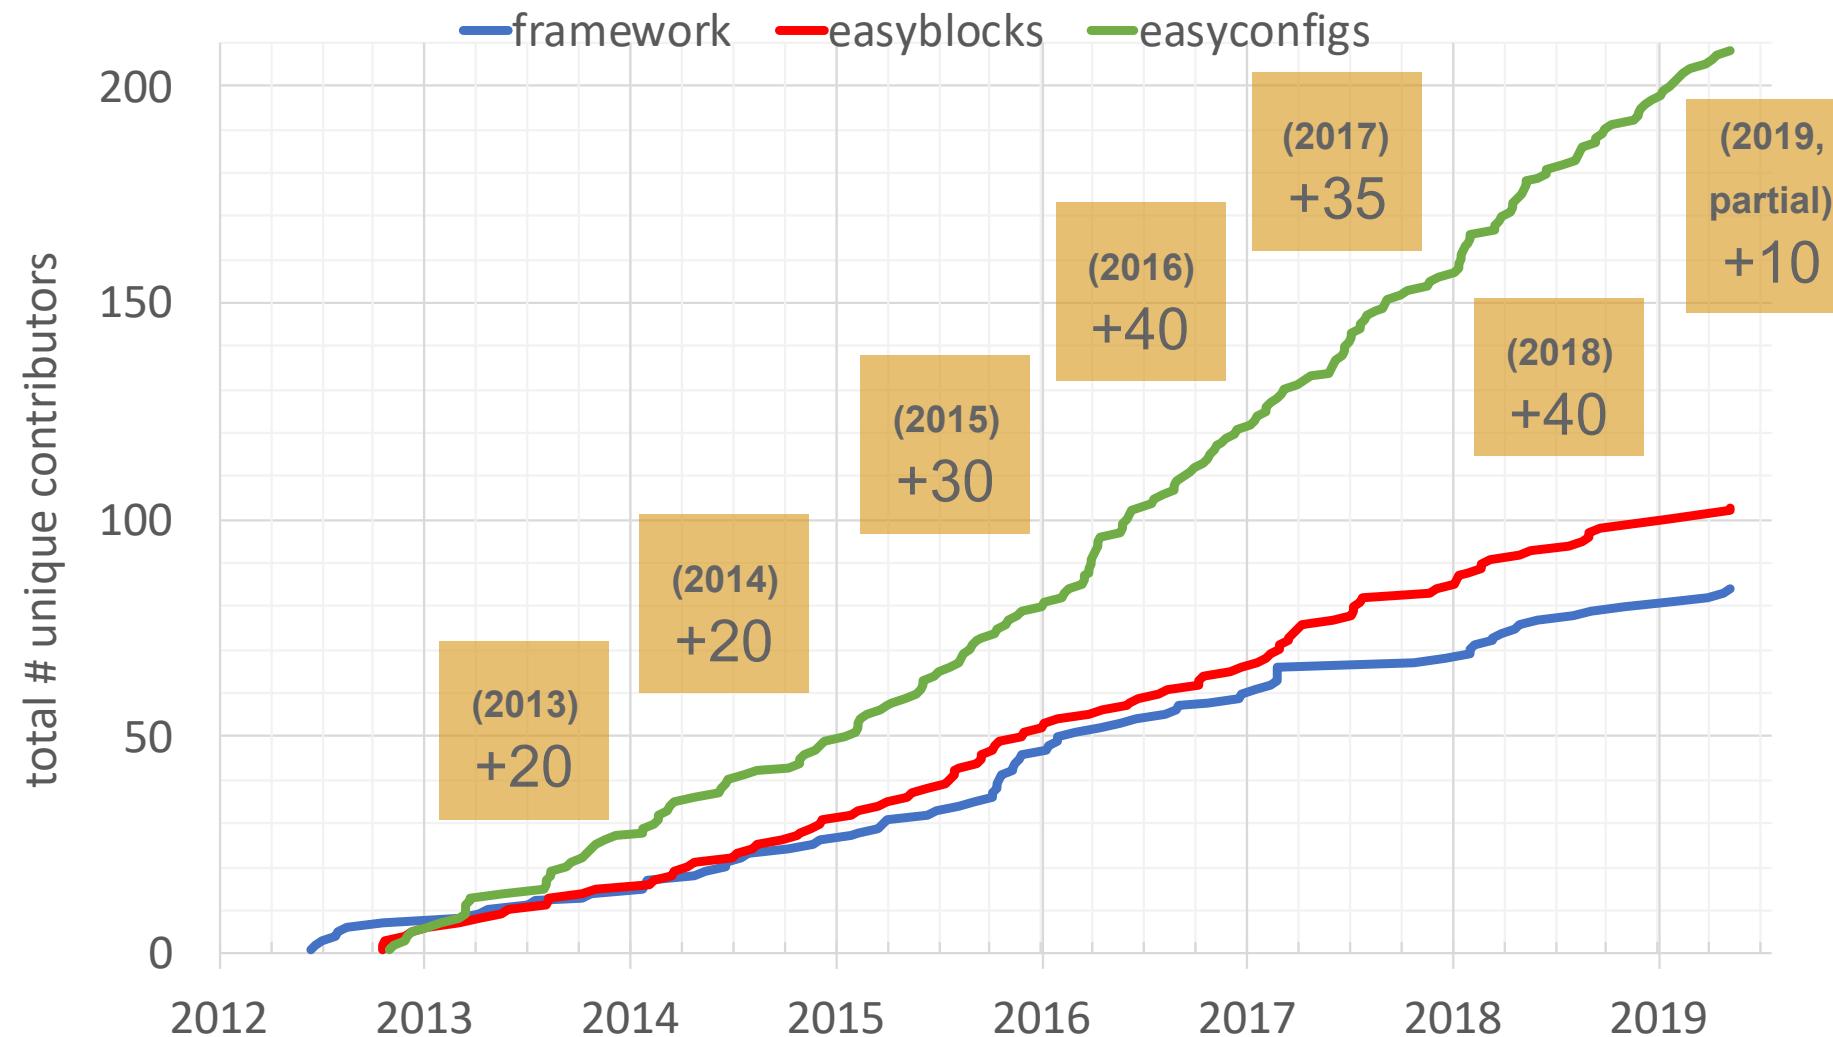
(source: Google Analytics)

The sun never sets on EasyBuild...



cities from where <https://easybuild.readthedocs.io> was visited at least 10 times during the last year
(source: Google Analytics)

Unique contributors (total, per repository)



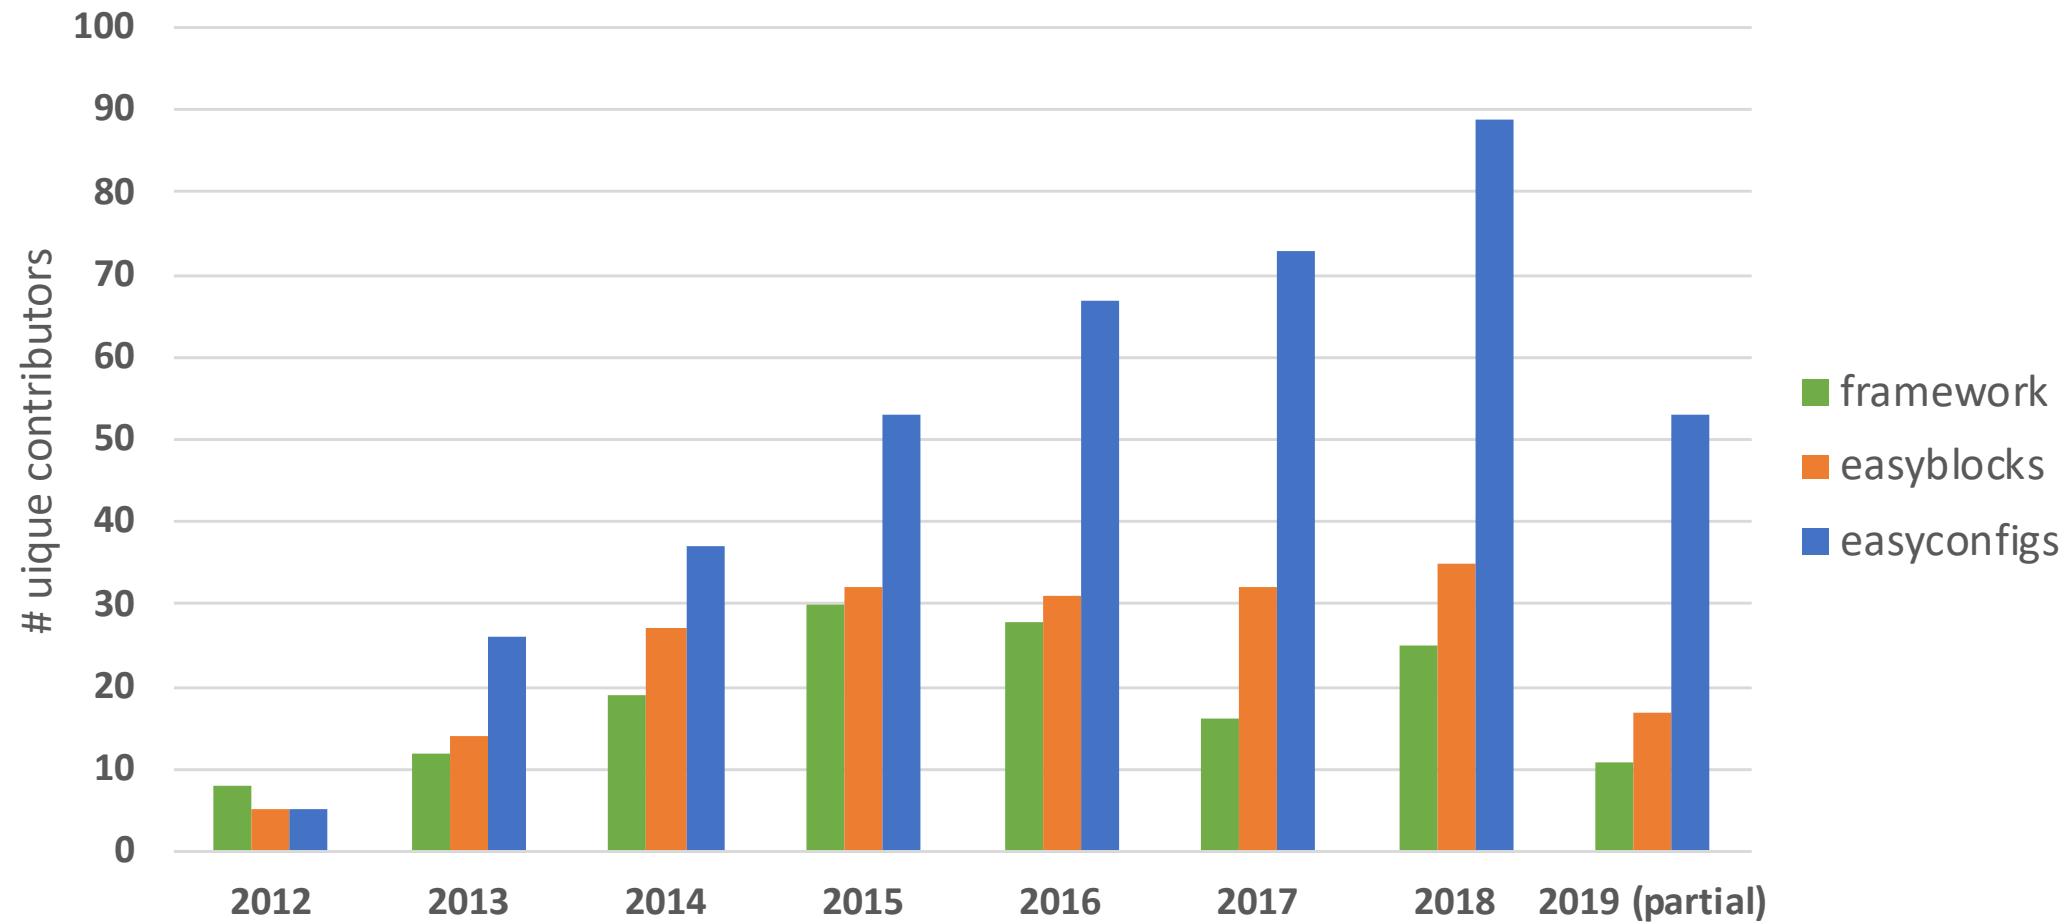
> 200 unique contributors for easyconfigs, > 100 for easyblocks, ~85 for framework

Unique contributors (per repository, per year)



Group of easyconfigs contributors is growing every year (reaching 100 in 2019?).

Stable group of easyblocks contributors, fluctuating a bit for framework.



Integration with Docker/Singularity



<https://easybuild.readthedocs.io/en/latest/Containers.html>

- (experimental) integration with Singularity since EasyBuild v3.6.0 (April 2018)
 - extended to also support Docker in EasyBuild v3.6.2 (July 2018)
- `eb --containerize` (or '`eb -C`') generates recipe to build container image
- generated recipe uses '`eb`' to build and install specified software in container image
- `--container-build-image` can be used to also build container image
 - makes `eb` call out to '`sudo singularity build`' (or '`sudo docker build`')
- **current implementation is experimental, has some rough edges...**

Integration with Docker/Singularity (demo)



<https://easybuild.readthedocs.io/en/latest/Containers.html>

```
$ export EASYBUILD_CONTAINER_BASE=shub:shahzebsiddiqui/eb-singularity:centos-7.4.1708
$ eb --experimental --containerize TensorFlow-1.13.1-foss-2018b-Python-3.6.6.eb
== temporary log file in case of crash /tmp/eb-sMtVBY/easybuild-1luKuM.log
== Singularity definition file created at /home/kehoste/containers/
Singularity.TensorFlow-1.13.1-foss-2018b-Python-3.6.6
== Temporary log file(s) /tmp/eb-sMtVBY/easybuild-1luKuM.log* have been removed.
== Temporary directory /tmp/eb-sMtVBY has been removed.
$ cat /home/kehoste/containers/Singularity.TensorFlow-1.13.1-foss-2018b-Python-3.6.6
Bootstrap: shub
From: shahzebsiddiqui/eb-singularity:centos-7.4.1708

%post
...
# change to 'easybuild' user
su - easybuild

eb TensorFlow-1.13.1-foss-2018b-Python-3.6.6.eb --robot --installpath=/app/ --prefix=/
scratch --tmpdir=/scratch/tmp

%environment
...
module load TensorFlow/1.13.1-foss-2018b-Python-3.6.6
```

Other features that were not covered in detail...



<http://easybuild.readthedocs.io>

- letting users manage their software stack on top of centrally provided modules
- installing hidden modules, hiding certain dependencies & toolchains
- support for using RPATH linking
- partial installations: only (re)generate module file, install additional extensions
- submitting installations as jobs to an HPC cluster via --job (distributed installation!)
- creating packages (RPMs, ...) for software installations done with EasyBuild
- installing (bundles of) Python packages for multiple Python versions in a single prefix

Outlook to easybuild v4.0

<https://github.com/easybuilders/easybuild/issues/447>

Ongoing effort, see '4.x' branches in
EasyBuild repositories on GitHub!

- ✓ support for running EasyBuild on top of Python 3.5 (or more recent Python 3)
- ✓ no more required dependencies (vsc-base has been ingested, setuptools no longer needed)
- WIP deprecated 'dummy' toolchain, replaced with 'system' toolchain (mostly about renaming)
 - custom easyblock for OpenMPI with more sensible default configuration
- WIP switch to using pip by default for installing Python packages (?)
 - (hopefully) stable integration with Singularity & Docker
 - **ETA: summer 2019**

- EasyBuild: GPLv2 license; Spack: MIT/Apache 2.0 license
- **no stable releases yet for Spack (< 1.0), EasyBuild is stable since 2012**
- on par w.r.t. amount of supported software (but differences w.r.t. which software)
- (originally) targeted for different use cases: HPC support teams vs developers
- fixed dependency/toolchain versions in EasyBuild vs flexible CLI in Spack
- running EasyBuild on top of Python 3 is not supported yet (but coming really soon!)
- macOS support in EasyBuild is limited (no toolchains/testing for macOS)
- both projects are backed by an active & supportive community!

For a more detailed comparison,
see https://archive.fosdem.org/2018/schedule/event/installing_software_for_scientists

Papers on easybuild

Modern Scientific Software Management Using EasyBuild and Lmod

Markus Geimer (JSC), Kenneth Hoste (HPC-UGent), Robert McLay (TACC)

http://easybuilders.github.io/easybuild/files/hust14_paper.pdf

Making Scientific Software Installation Reproducible On Cray Systems Using EasyBuild

Petar Forai (IMP), Guilherme Peretti-Pezzi (CSCS), Kenneth Hoste (HPC-UGent)

https://cug.org/proceedings/cug2016_proceedings/includes/files/pap145.pdf

Scientific Software Management in Real Life: Deployment of EasyBuild on a Large Scale System

Damian Alvarez, Alan O'Cais, Markus Geimer (JSC), Kenneth Hoste (HPC-UGent)

<http://easybuilders.github.io/easybuild/files/eb-jsc-hust16.pdf>



Questions?

Kenneth Hoste (HPC-UGent)

kenneth.hoste@ugent.be

20190510 - CSCS Software Management Course

https://users.ugent.be/~kehoste/EasyBuild_20190510_CSCS_software_management.pdf

<https://easybuilders.github.io/easybuild>

<https://easybuild.readthedocs.io>

Extra slides

(only used in case of time/interest)

Using a custom module naming scheme



- a couple of different module naming schemes are included in EasyBuild
 - see `--avail-module-naming-schemes`
 - specify active module naming scheme via `--module-naming-scheme`
 - default: `EasyBuildMNS (<name>/<version>-<toolchain>-<versionsuffix>)`
- **you can implement your own module naming scheme relatively easily**
 - specify how to compose module name using provided metadata
 - via Python module that defines custom derivative class of `ModuleNamingScheme`
 - make EasyBuild aware of it via `--include-module-naming-schemes`
- decouple naming of install dirs vs modules via `--fixed-installdir-naming-scheme`

Flat module naming scheme



- all modules are always available for loading
- long(er) module names
- 'module avail' may be overwhelming for users
- too easy to load incompatible modules together

legend

(not available)
(available)
(loaded)

GCC/5.3.0

GCC/6.1.0

OpenMPI/1.10.2-GCC-5.3.0

OpenMPI/2.1.0-GCC-5.3.0

OpenMPI/1.10.3-GCC-6.1.0

OpenMPI/2.1.0-GCC-6.1.0

FFTW/3.3.4-gompi-2016.04

FFTW/3.3.6-gompi-2016.04

FFTW/3.3.4-gompi-2016.07

FFTW/3.3.6-gompi-2016.07

Hierarchical module naming scheme (1)

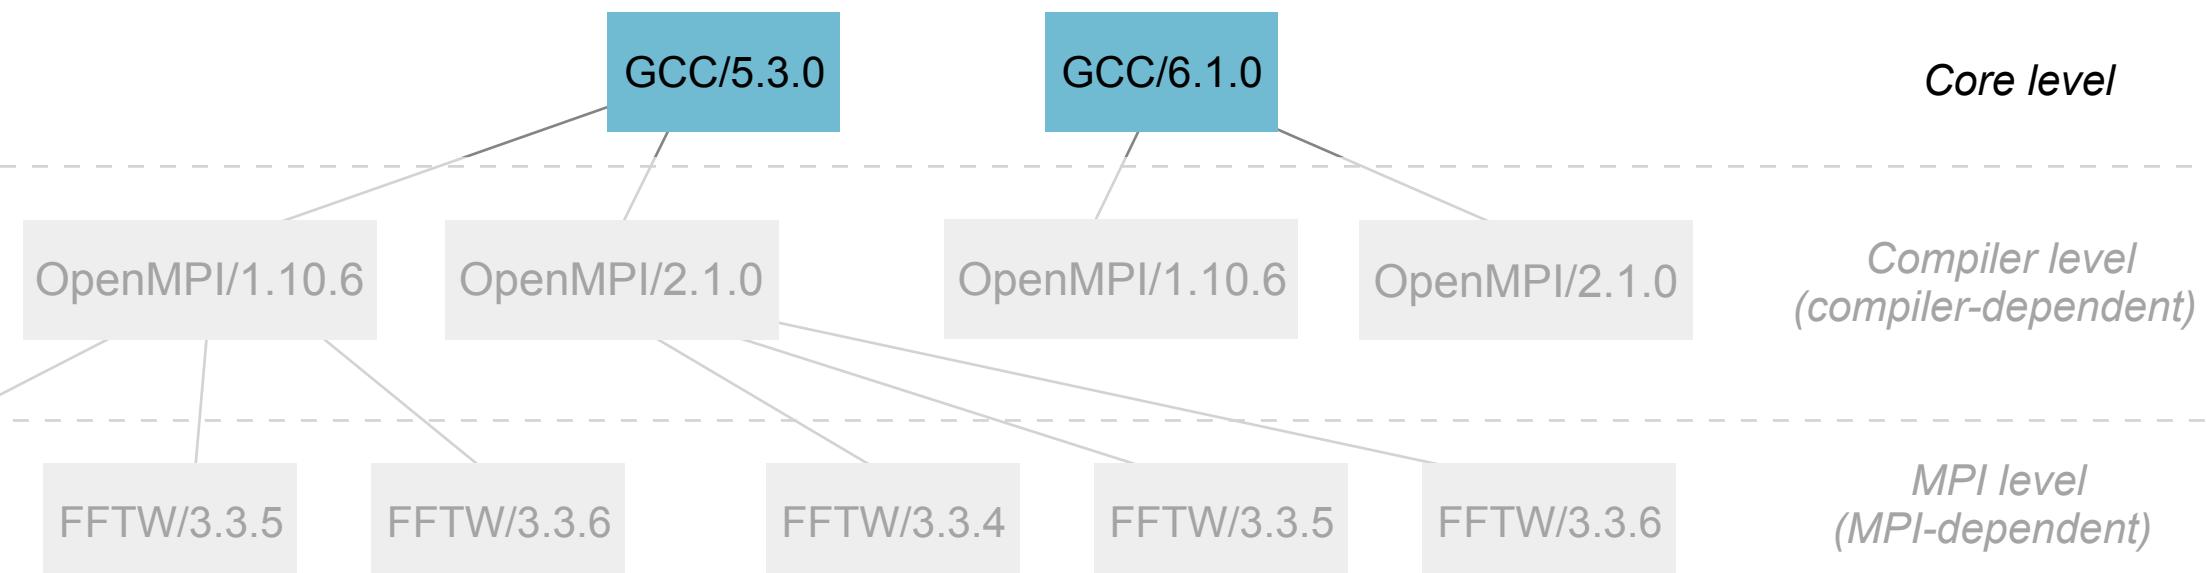


- modules are organised in a tree-like fashion
- initially, only 'core' modules are available for loading
- typically 3 hierarchy levels: core, compiler-dependent, MPI-dependent
- other modules are only visible via 'module spider'

legend

(not available)
(available)
(loaded)

Core level



*MPI level
(MPI-dependent)*

*Compiler level
(compiler-dependent)*

Core level

Hierarchical module naming scheme (2)



- Core modules may extend \$MODULEPATH with an additional location
- loading a Core module may make more modules available
- in this example, loading a GCC module makes OpenMPI modules available

legend

(not available)

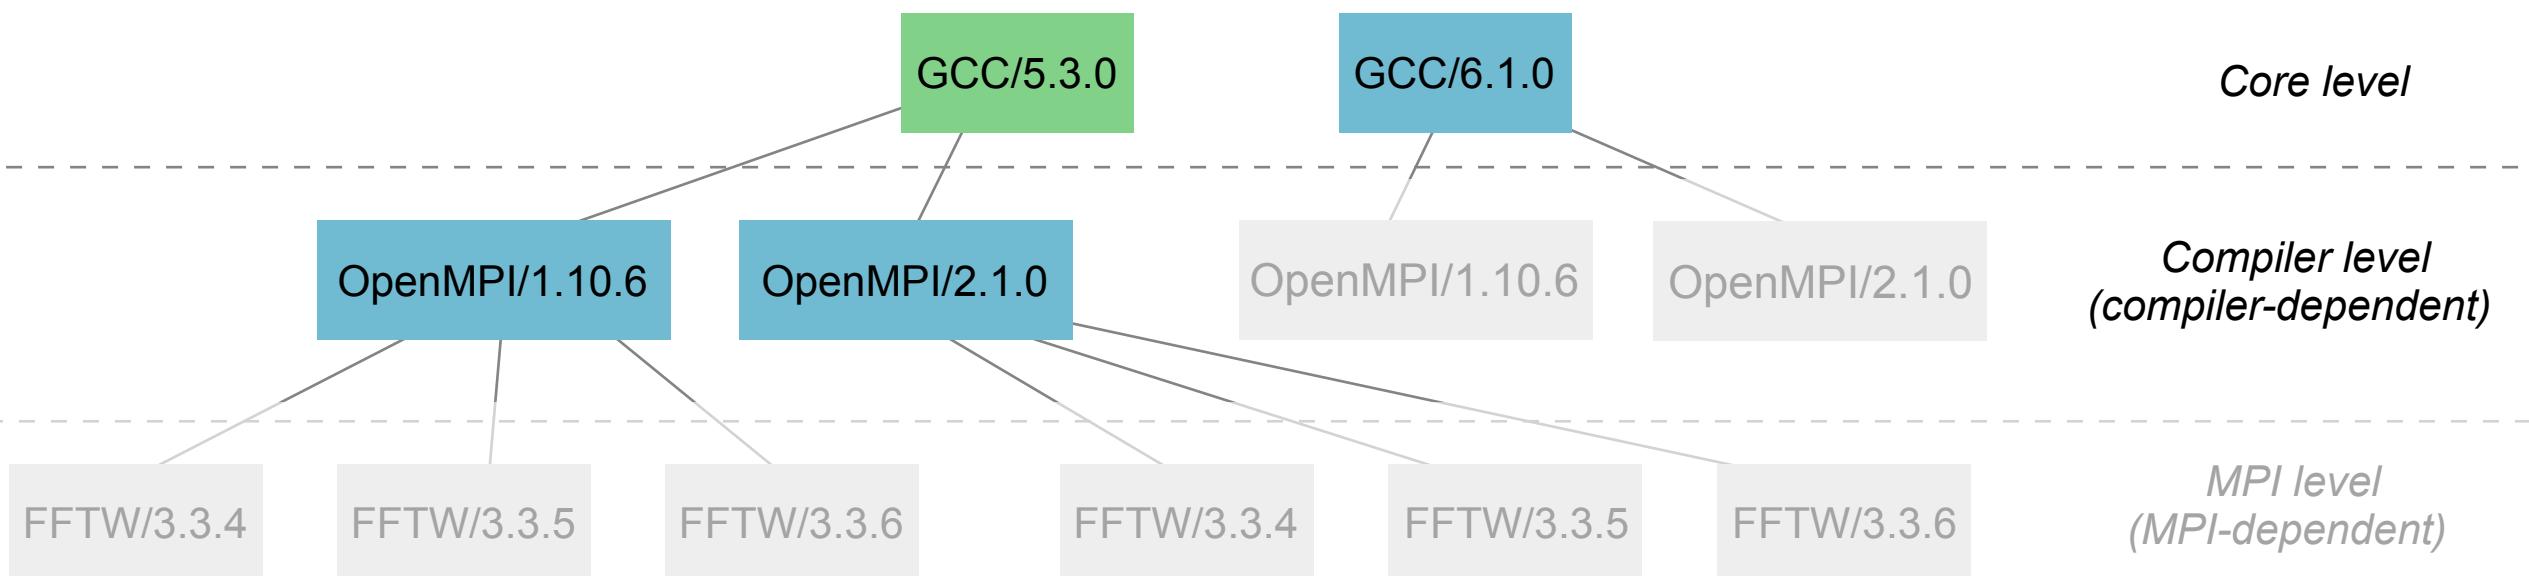
(available)

(loaded)

Core level

Compiler level
(compiler-dependent)

MPI level
(MPI-dependent)



Hierarchical module naming scheme (3)



- even more modules may be made available by loading other modules
- for example, loading an OpenMPI modules reveals MPI-dependent modules
- EasyBuild can organise modules in hierarchy for you!

legend

(not available)
(available)
(loaded)

Core level

